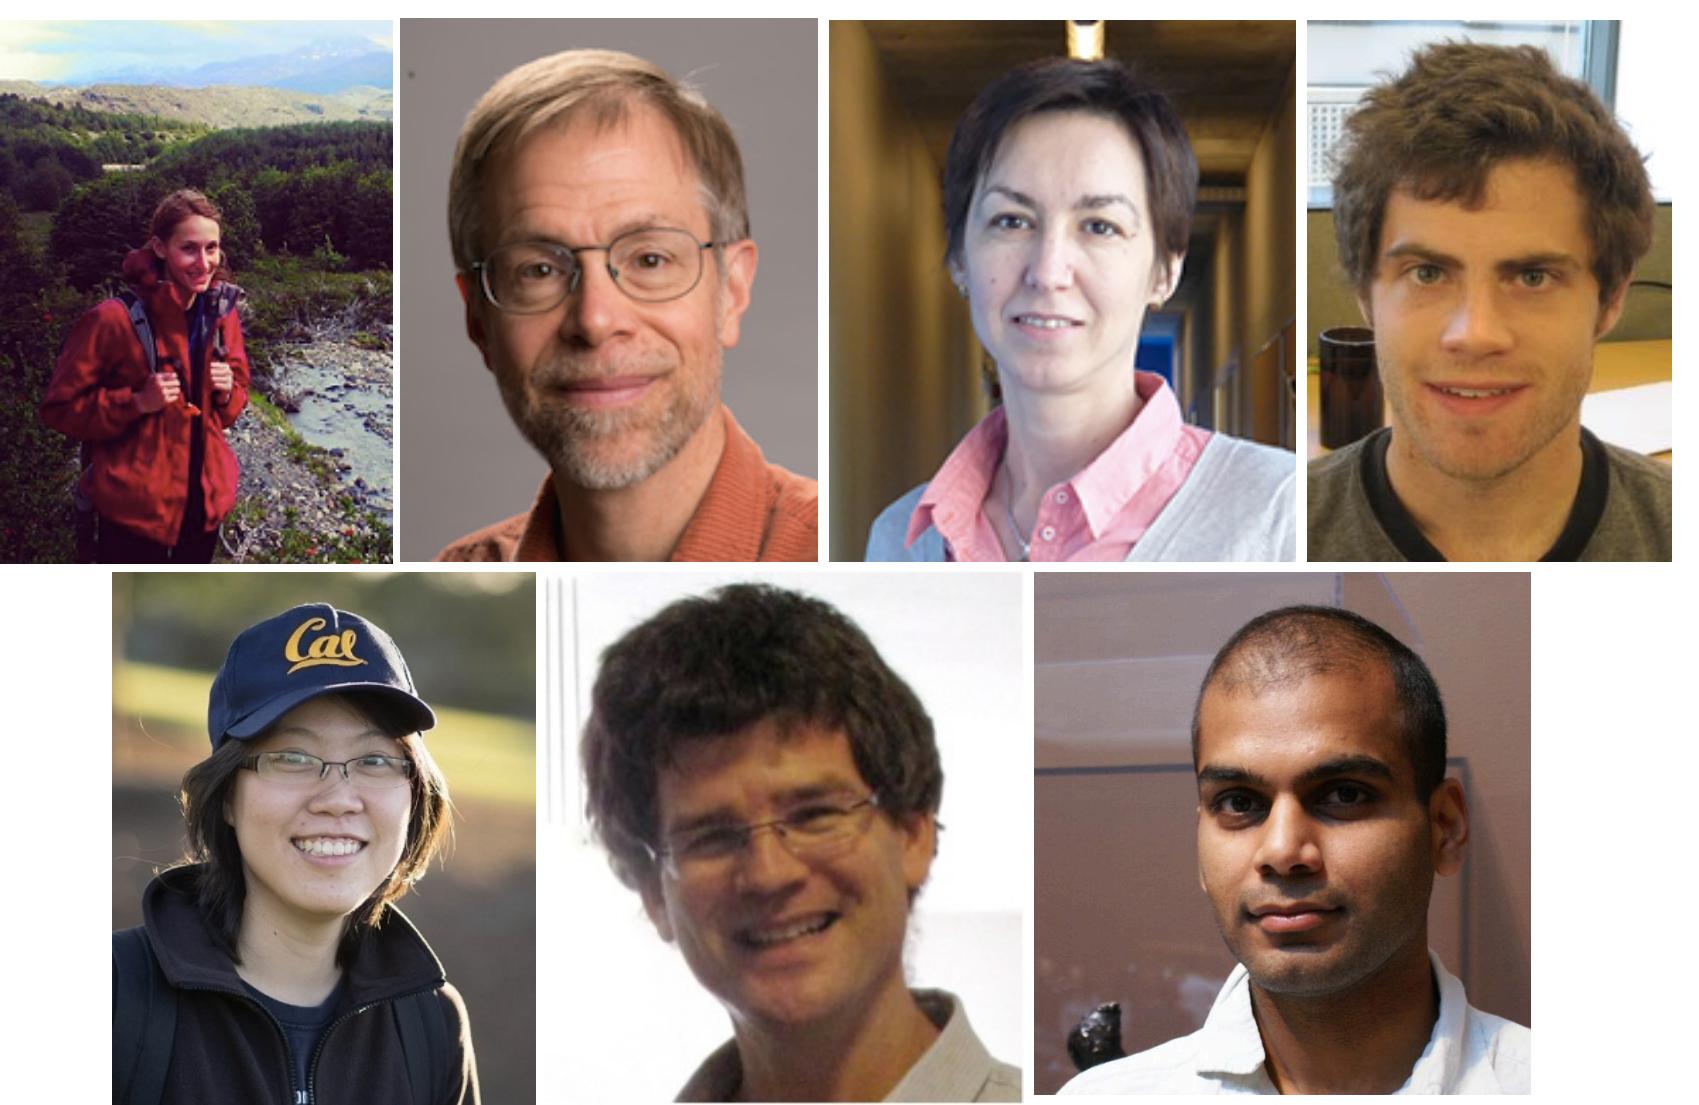


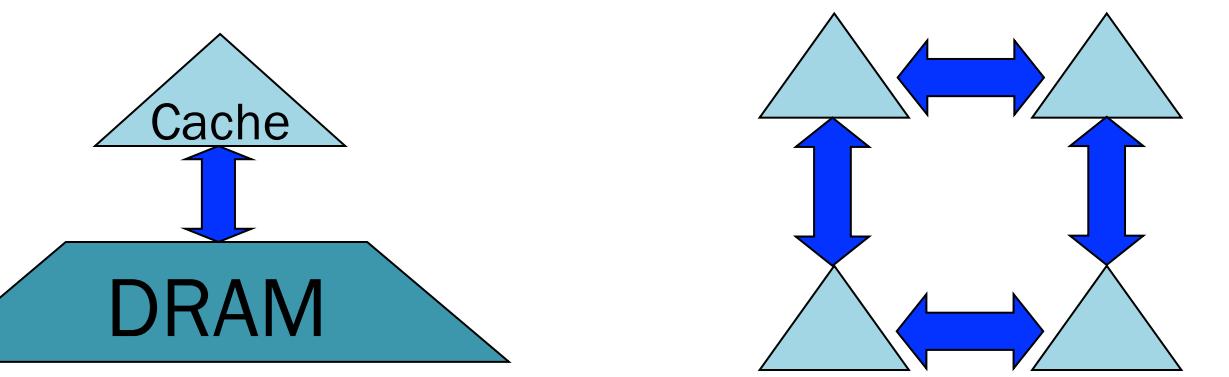
Write-Avoiding Algorithms

Erin Carson, James Demmel, Laura Grigori, Nicholas Knight,
 Penporn Koanantakool, Oded Schwartz, Harsha Vardhan Simhadri

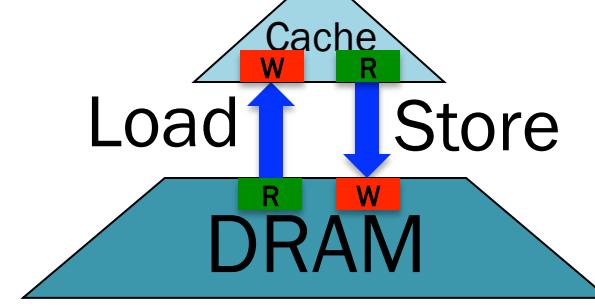


Motivation / Model

- Moving data (communication) most expensive operation (in time or energy), so avoid it
 - Work so far on Communication-Avoiding Algorithms
 - Provably minimize #loads and #stores between levels of memory hierarchy, and #words sent over network
 - big speedups in theory and practice



- Loads and Stores are not equal
 - Load = read from slow mem, write to fast mem
 - Store = read from fast mem, write to slow mem



- Writes may be much more expensive than reads:
 - Nonvolatile Memory (NVM)
 - Flash, PCM, STT-RAM, ReRAM, ...
 - STT-RAM $\text{read} = .14\text{ns} / 10^{-15}$ Joules vs. $\text{write} = 10\text{ns} / 10^{-12}$
 - Faster wear-out with writes than reads
 - When intermediate results written to disk for fault tolerance (e.g. Spark)
 - Writes may cause more coherency traffic in a shared memory environment

- Can we minimize #writes, not just #loads+#stores?

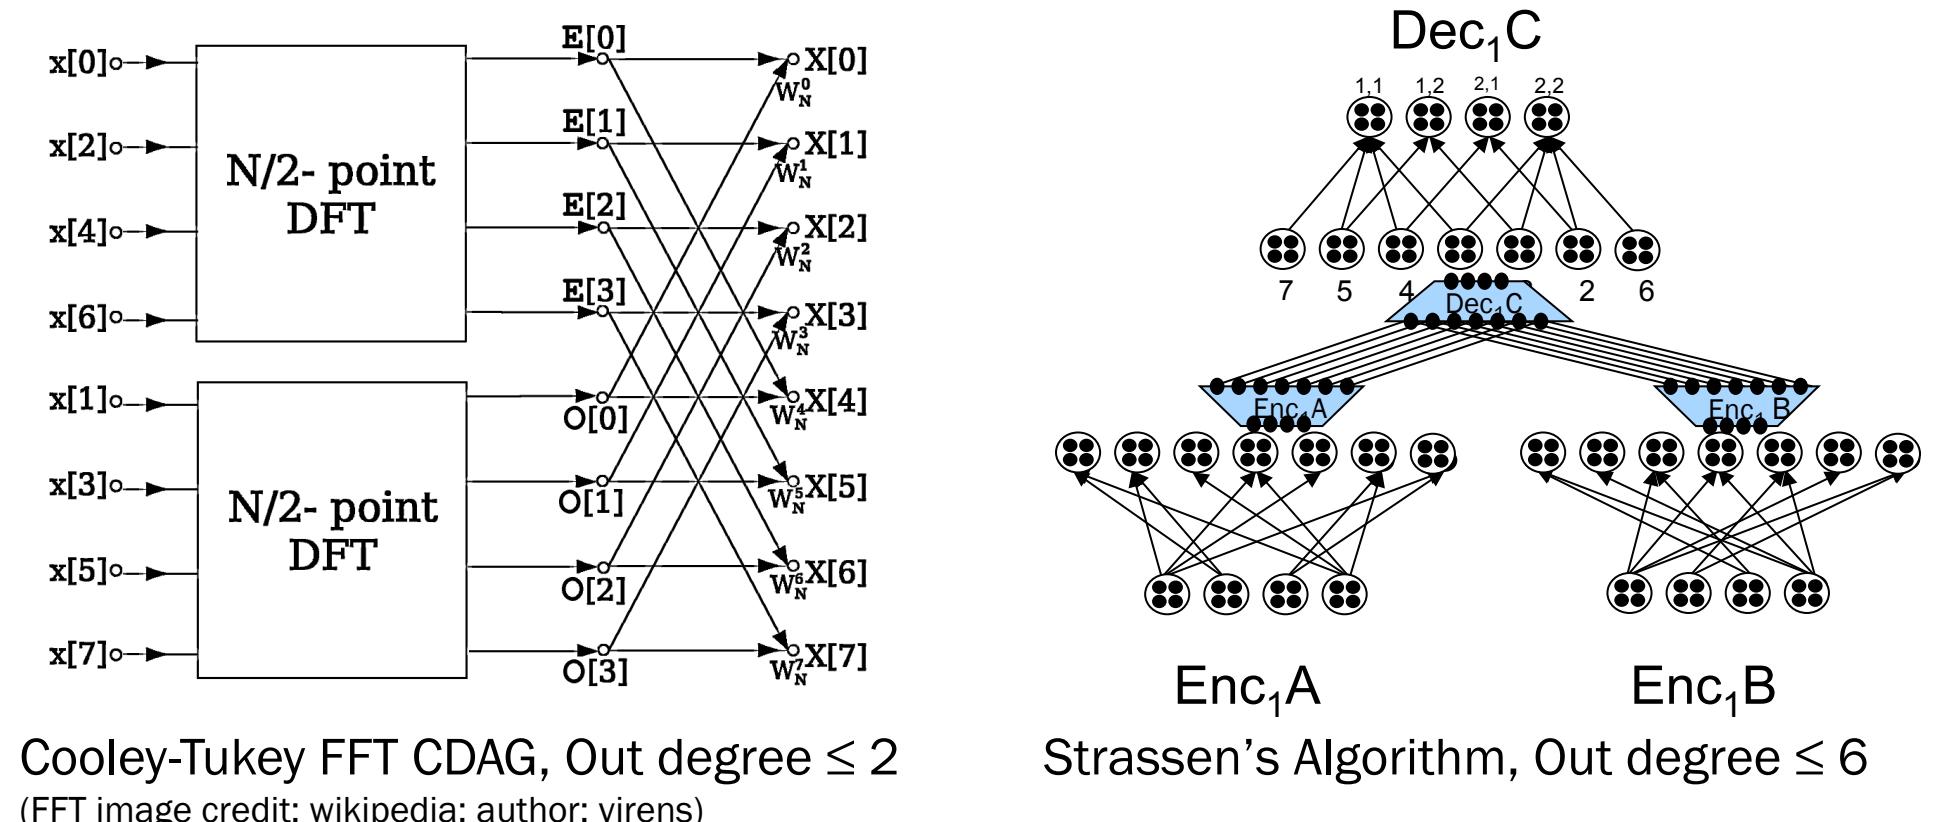
- Theorem 1:**
 - #writes to fast memory $\geq \frac{1}{2} (\#loads + \#stores)$
 - #writes to slow mem \geq size of output

Take-away: may be able to do fewer writes to slow memory!
 When is this (much smaller) lower bound attainable?

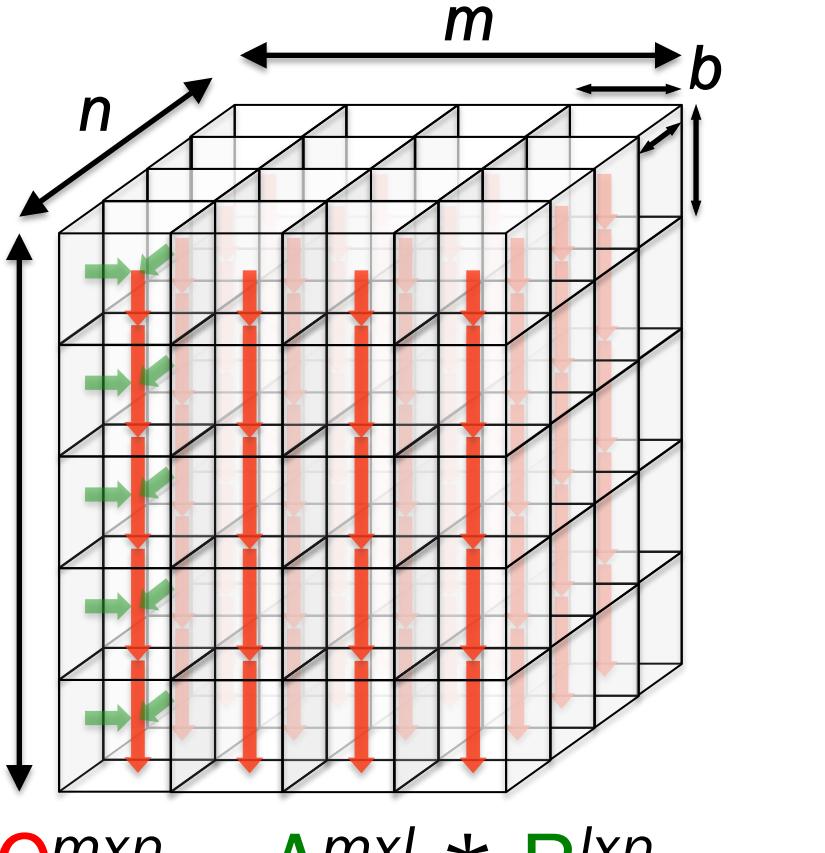
- Definition: An algorithm that is CA and also attains an asymptotically smaller lower bound on #writes to slow memory is called Write Avoiding (W-A).

Write Lower Bounds

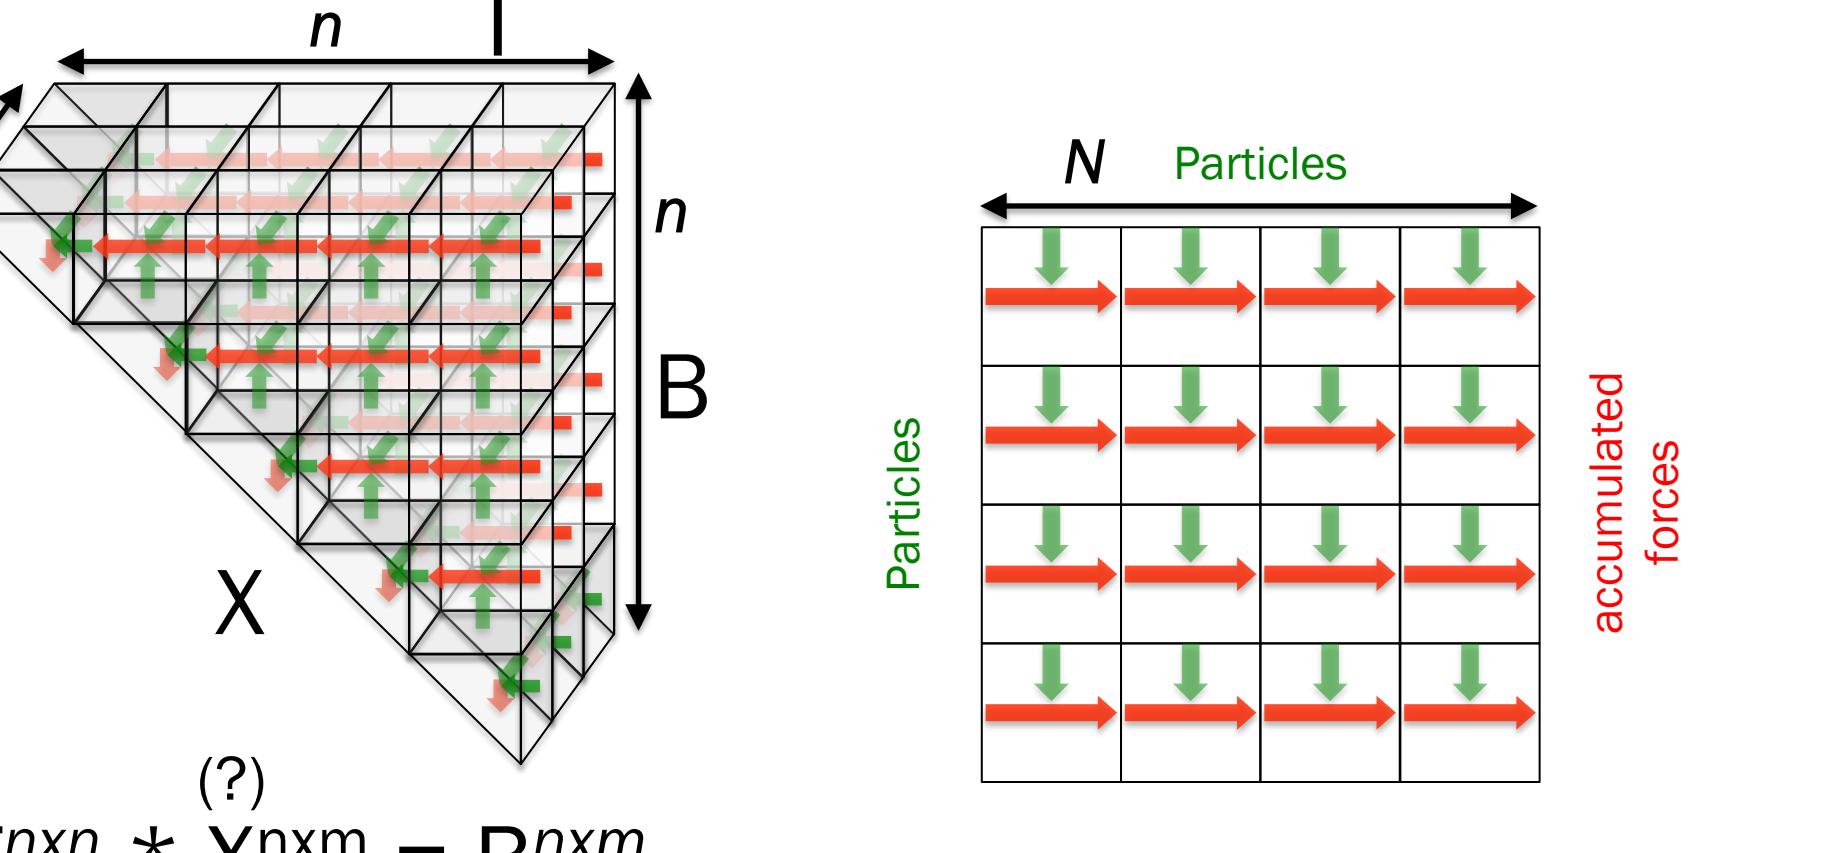
- W-A algorithms don't always exist
- CDAG of an algorithm and its input is a directed graph
 - vertices = arguments (inputs, outputs, intermediate data)
 - edges = direct dependencies
- Theorem 2:** If the out-degree of a CDAG of an algorithm (or large portions of it) is bounded by d , then
 - #writes to slow memory $\geq (1/d) * \#reads$ from slow memory
 - [see paper for a more precise statement]
 - Corollary : Algorithms with bounded out-degree CDAGs like Cooley-Tukey FFT ($d \leq 2$) and Strassen's Algorithms ($d \leq 6$) cannot be made write-avoiding by instruction ordering.



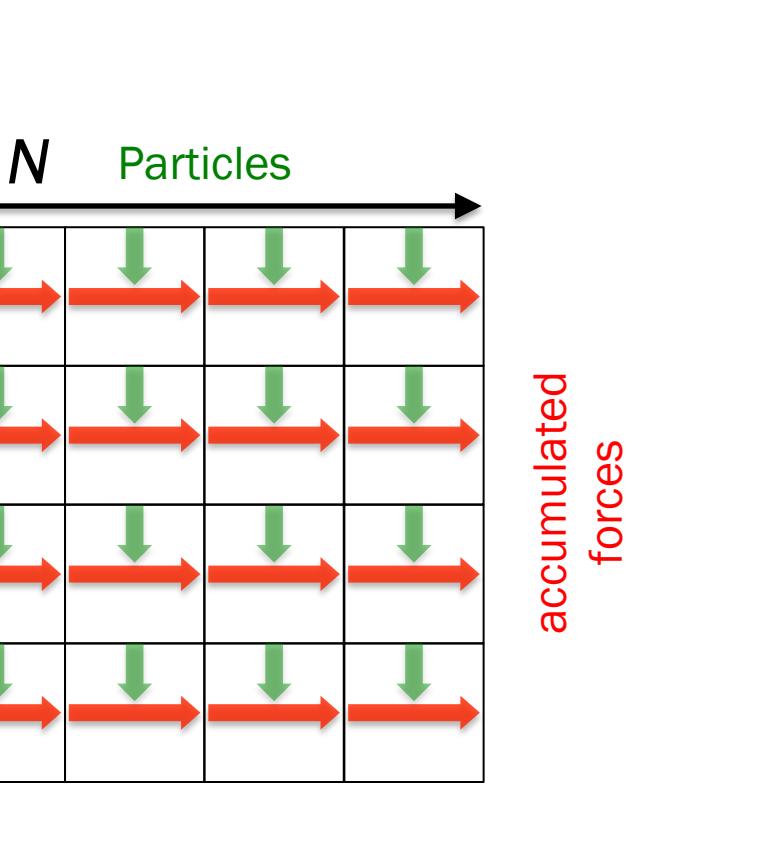
Write-Avoiding Linear Algebra and N-Body Algorithms



Classical Matrix Multiplication



Triangular Solve



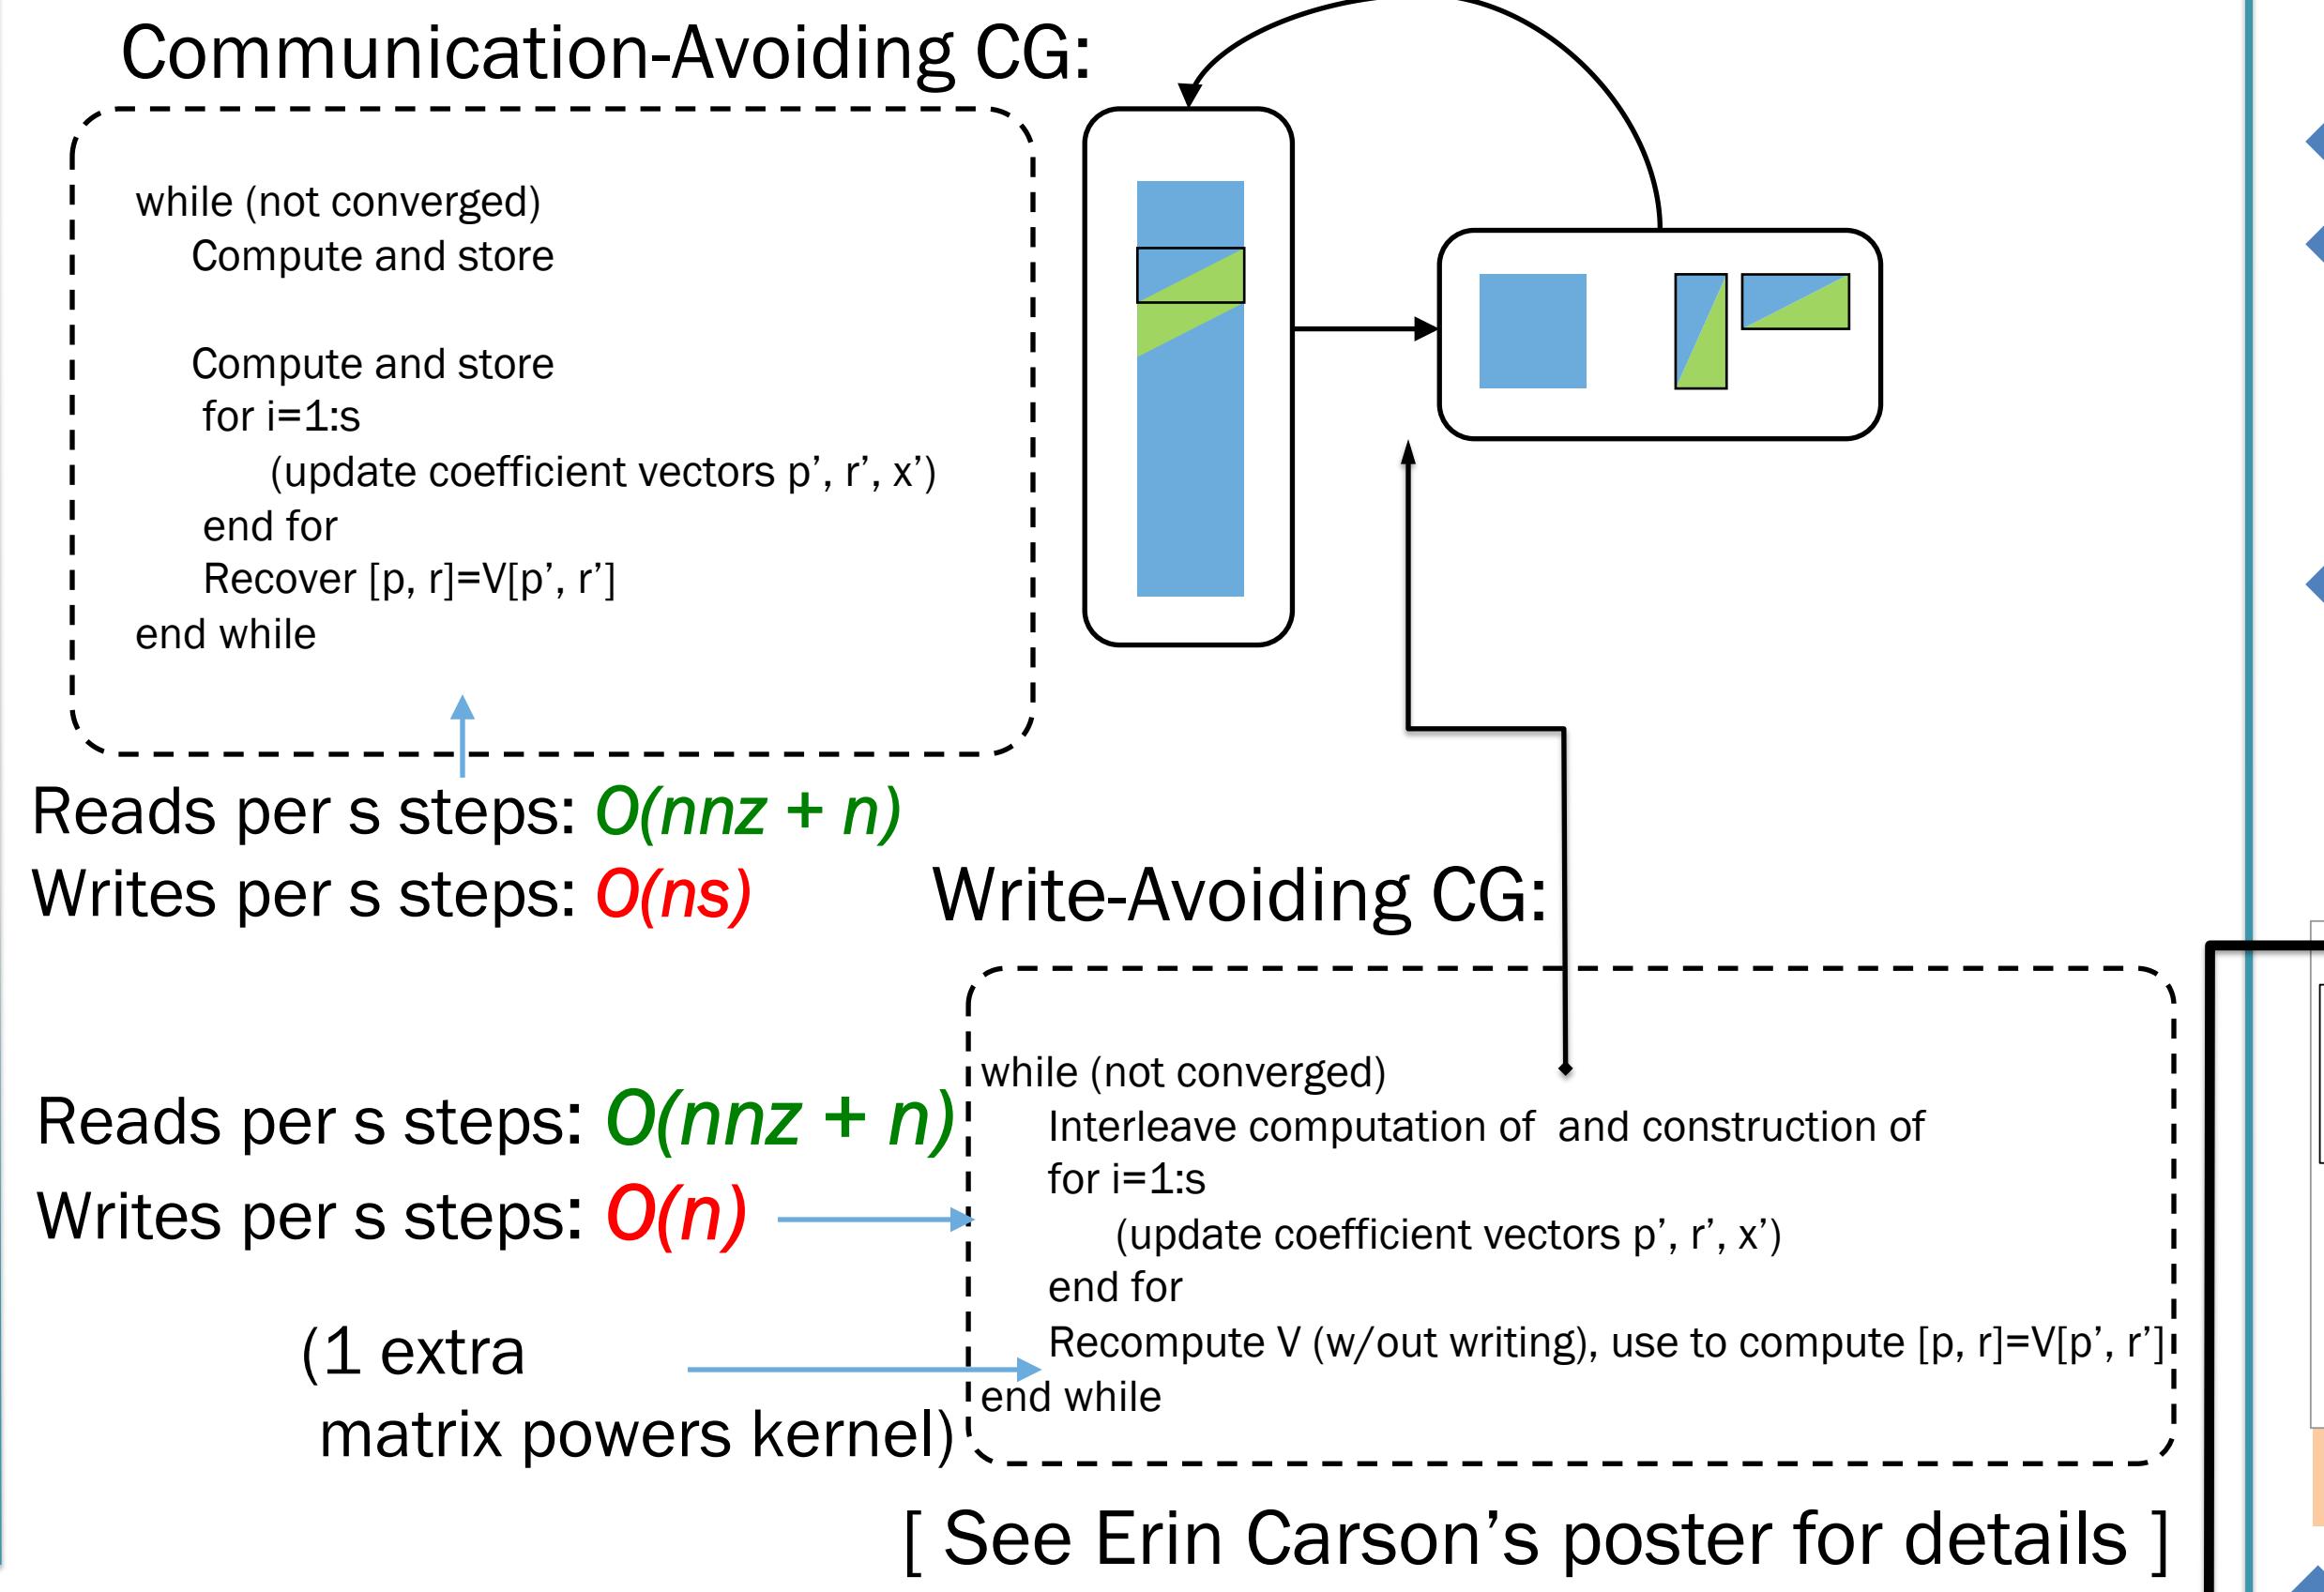
Direct N-body

Writes to fast memory:
 $ml + 2mnl/b$
 Stores to slow memory:
 mn

Loads to fast memory:
 $mn + n^2/2 + mn^2/b$
 Stores to slow memory:
 N

- There are sequential W-A algorithms for classical matrix multiplication, Triangular Solve, Cholesky factorization and direct N-body algos.
- Based on appropriate loop reordering
- All use
 - Explicit blocking based on cache size
 - extend to multiple levels of memory
 - assume explicit control over data movement
 - special cases of CA algorithms
- Not all communication avoiding algorithms are W-A
- Conjecture: should work for many other familiar algorithms

Krylov Subspace Methods

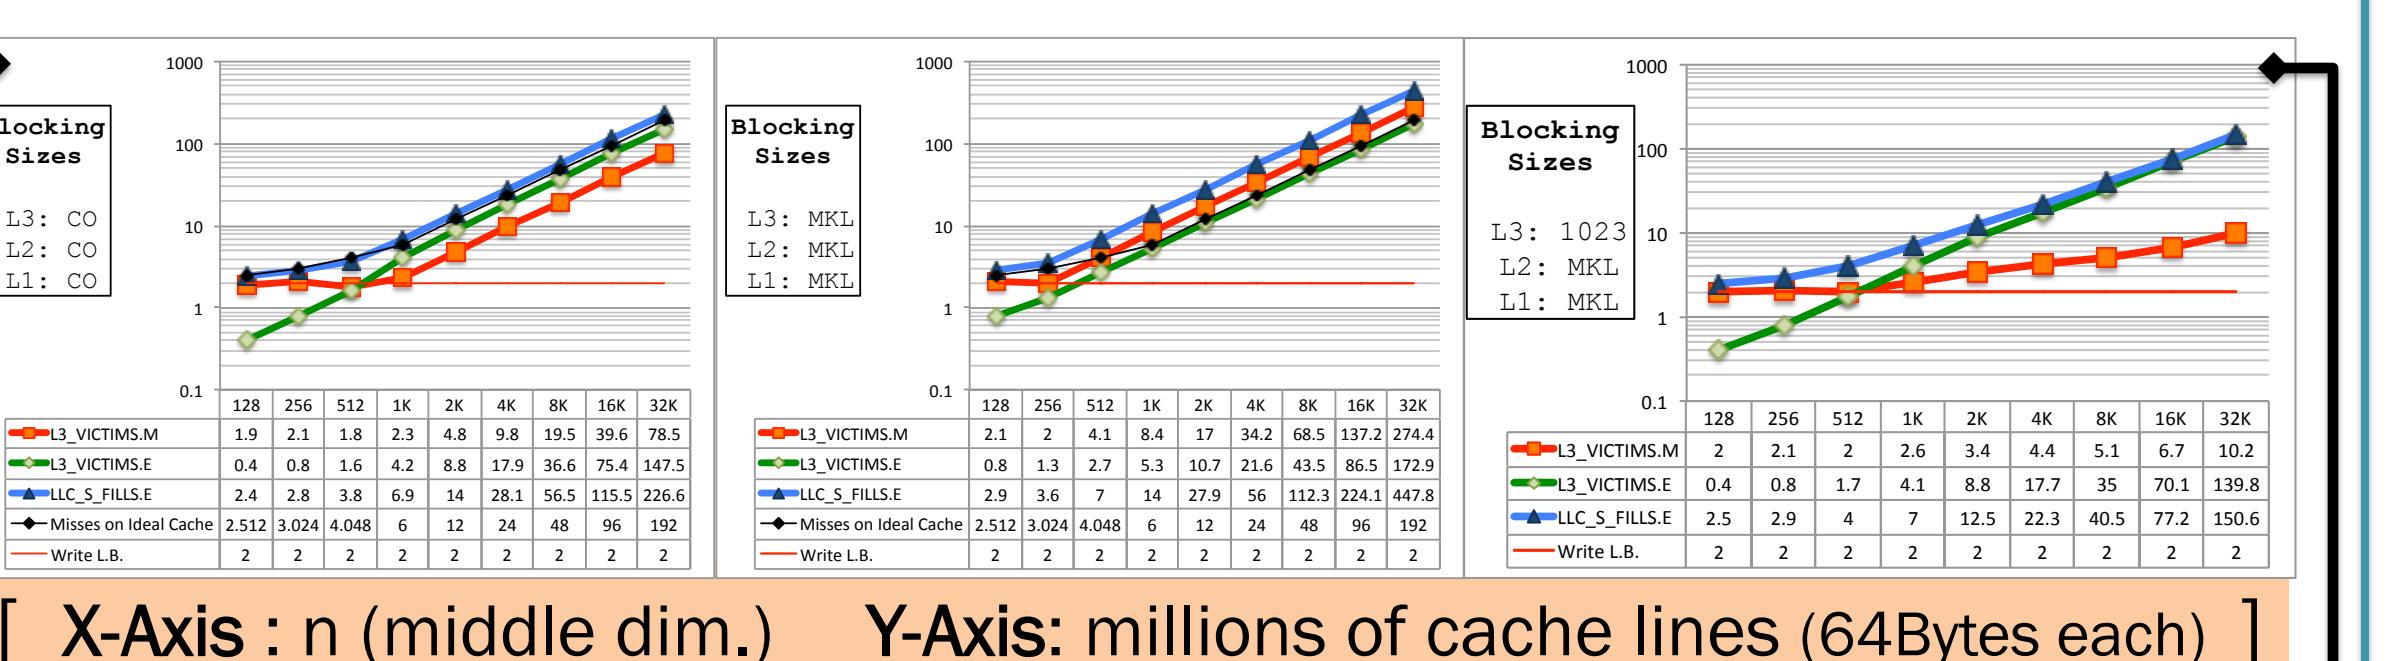


Cache-Oblivious → Not W-A

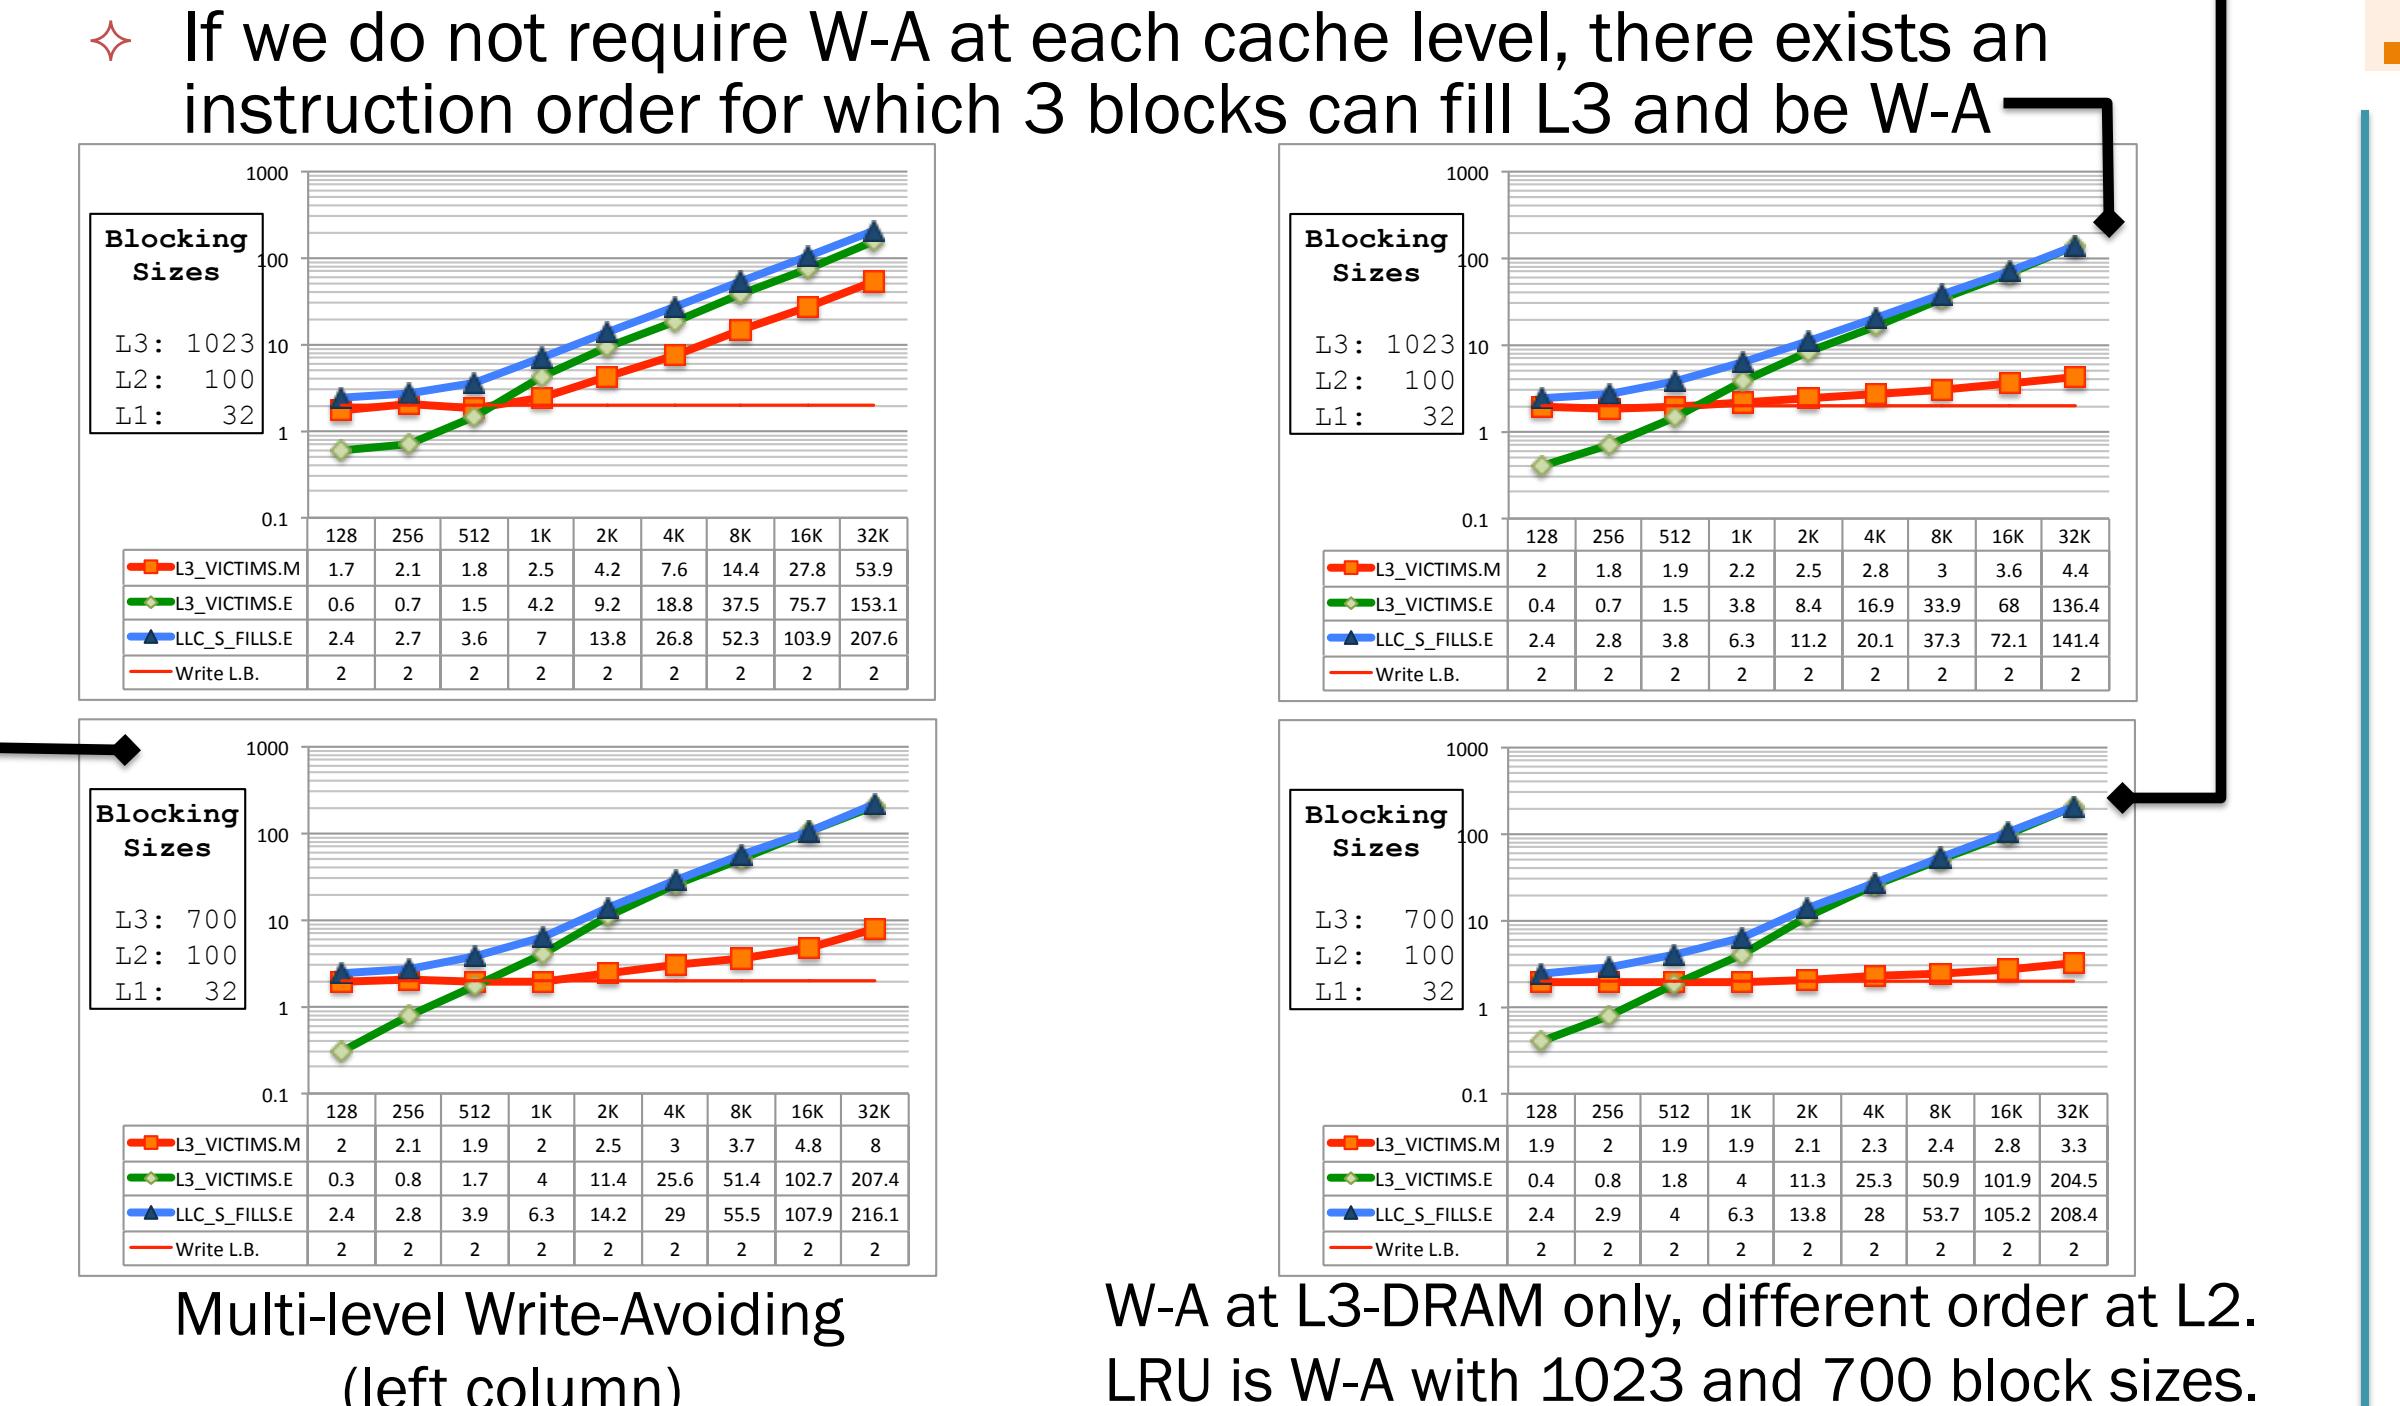
- Definition: A Cache-Oblivious (C-O) algorithm does not depend on cache-size parameters
 - There are well-known sequential Cache-Oblivious algorithms for classical matrix multiplication, TRSM, Cholesky, etc. that are communication avoiding for general nested memory hierarchies
- Can CO algorithms be Write-Avoiding?
 - No!
- Theorem 3: For a large class of problems, Cache-Oblivious algorithms do at least a constant fraction as many writes to slow memory as reads from slow memory
 - "Large class" means "smells like 3-nested loops"
- Conjecture: extends to nested loops accessing arrays
- Empirical measurements of C-O Matrix Multiplication are in line with Theorem 3

Hardware Perf Counters & Cache Replacement Policy

- W-A algorithms above assume explicit control over data movement (Ex: User accesses Flash over PCIe)
- Do cache policies like LRU enable WA algorithms?
- Experimental Data for 4000 x n x 4000 matrix mult
 - Intel Nehalem-EX Xeon 7560, MESIF coherence, pseudo-LRU L3
 - 24MB L3, 256KB L2, 32KB L1 (Output: 122MB = 2.0mln cache lines)
 - hugectl used to allocate huge pages, avoid TLB misses
- C-box events accessed via customized Intel PCM 2.4
 - L3_VICTIMS.M = #stores from L3 to DRAM (Modified L3 evictions)
 - L3_VICTIMS.E = #L3 evictions without DRAM writes (Exclusive lines)
 - LLC_S_FILLS.E = #loads from DRAM to L3 in Exclusive state

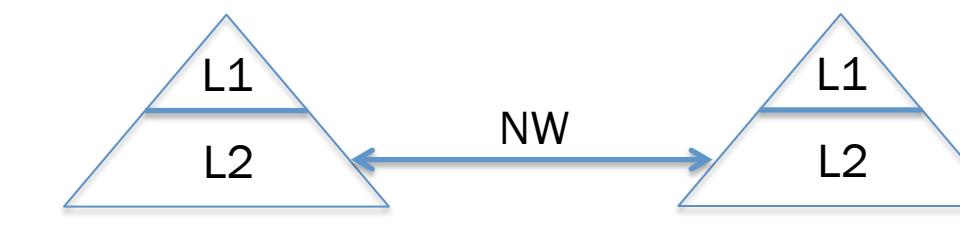


- Measurements for Cache-Oblivious, MKL and W-A
 - Cache-Oblivious optimizes for reads, but not writes to DRAM
 - MKL optimizes for neither (better time though)
 - W-A at L3-DRAM (with $b=1023$, 3 blocks fit in L3) + MKL for L1 and L2 does fewer writes, but still not very close to optimal. Gap between lower bound and writes due to replacement policy.
- Proposition 6.1:** If 5 blocks fit in L3 cache (not just 3 blocks required in explicit blocking), LRU matches lower bound for any instruction order of block multiplication within L3 cache.
 - If we do not require W-A at each cache level, there exists an instruction order for which 3 blocks can fill L3 and be W-A

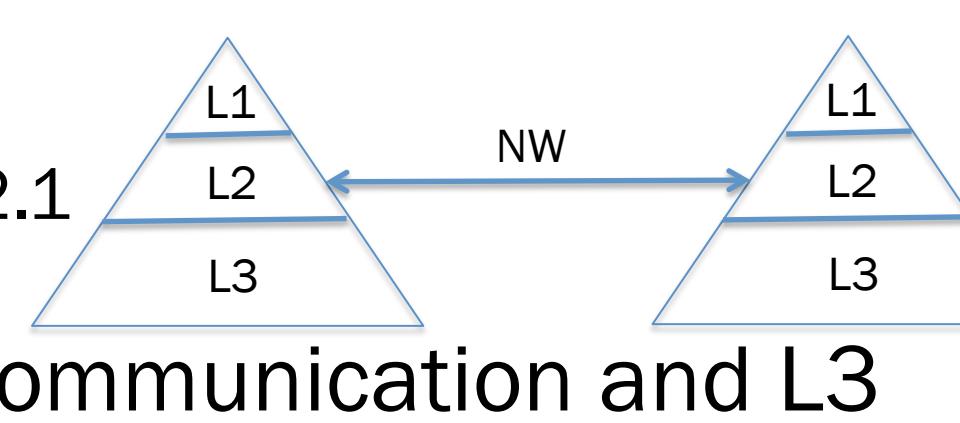


Parallel W-A Algorithms

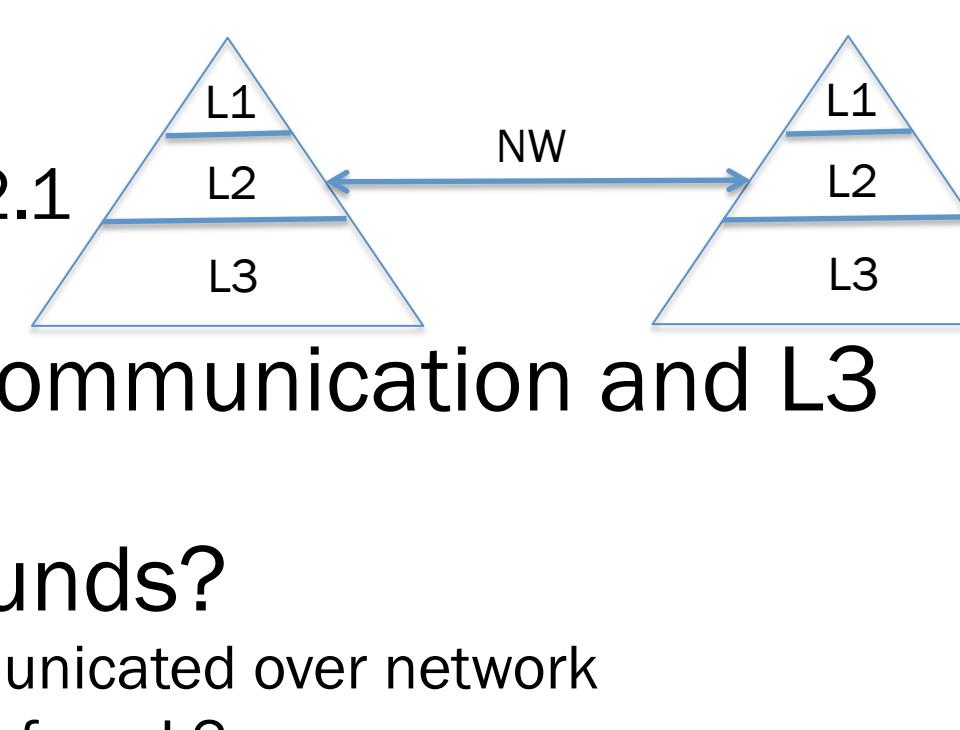
- Model 1:**
 - $L1$ = cache
 - $L2$ = DRAM
 - Network (NW) connects L2s
- Can we minimize NW communication and L2 writes from L1?
- Natural idea:
 - Use CA algorithm to minimize NW communication
 - Use WA algorithm locally on each processor
 - Applies to Matmul, TRSM, Cholesky, N-Body, ...
- Does it work for Matmul?
 - Goals: $n^2/P^{1/2}$ words moved on NW, n^2/P writes to $L2 \leftarrow L1$
 - Yes for NW, but $n^2/P^{1/2}$ writes to L2 from L1
 - Probably OK, since dominated by NW costs
 - Can attain both lower bounds, but with $P^{1/2}$ times as much L2, probably not worth it



- Model 2.1:**
 - $L3$ = NVM
 - Data fits in L2
- Is it worth using L3?
- Idea: use 2.5D algorithms that replicate data to reduce NW traffic
 - Ex: 2.5DMM moves $n^2/(cP)^{1/2}$ words over NW if enough memory for c copies of data ($1 \leq c \leq P^{1/3}$)
 - Using L3 may let us increase c , at cost of L3 writes
- Performance model
 - c_2 = #copies using L2, c_3 = #copies using L3, so $c_3 > c_2$
 - β_{NW} = network BW, β_{23} = L3 write BW, β_{32} = L3 read BW
 - Potential speedup = $(c_3/c_2)^{1/2} [\beta_{NW} / (\beta_{NW} + 1.5\beta_{23} + \beta_{32})]$



- Model 2.2**
 - Same architecture as Model 2.1
 - Data fits in L3, not L2
- Can we minimize NW communication and L3 writes from L2?
- Can we attain all lower bounds?
 - $W_{NW} = \Omega(n^2/(cP)^{1/2})$ words communicated over network
 - $W_{23} = \Omega(n^2/P)$ words written to L3 from L2
- Theorem 4 (bad news):** It is impossible to attain both lower bounds (See paper for details)
- Good news: There are algorithms that can attain either bound (but not the other)
 - Alg 1: $W_{NW} = \Omega(n^2/(cP)^{1/2})$ but $W_{23} = W_{NW}$
 - Alg 2: $W_{23} = \Omega(n^2/P)$ but $W_{NW} = \Omega(n^3/(P \cdot \text{size_of_L2})^{1/2})$
 - Which one is best depends on algorithmic & HW parameters
 - Extends to LU



- Possible to asymptotically reduce #writes to lowest level of memory hierarchy for many algorithms
 - Enables saving time and energy for nonvolatile memory
 - Works for many of the 7 dwarfs: Dense/Sparse LA, Un/structured grids, N-body, (not FFT)
- Future Work:**
 - Extend WA approach to other algorithms
 - Conjecture: extends to nested loops accessing arrays (HBL)
 - Conjecture: For direct N-body, need to double flops
 - Conjecture: W-A impossible for $O(n \log n)$ sorting or DFT
 - Can we minimize #messages too?
 - Extend theory, scheduling algorithms to shared memory
 - Implement on Firebox++
 - When does 'single node + NVM' beat 'cluster + DRAM'?